

目录

Abstract	5
List of Figures	6
List of Tables	7
1 Introduction	8
1.1 Context and Motivation	8
1.2 Aims and Objectives	8
1.3 Report Structure	8
2 Background	9
2.1 The Foreign Exchange Market	10
2.2 Time Series Forecasting	10
2.2.1 Classification	10
2.2.2 Regression	10
2.3 Currency Exchange Rate Forecasting	10
2.3.1 Statistics Approach	10
2.3.2 Classical Machine Learning Approach	10
2.3.3 Deep Learning Approach	10
2.3.4 Hybrid Methods	10
2.4 Hyperparameter Optimization	10

2.5	Critical Analysis	10
3	Methodology	11
3.1	Data Collection	11
3.1.1	Primary Dataset	12
3.1.2	Supplementary Features for Multivariate Models	12
3.1.3	Benchmarking Datasets	12
3.2	Data Analysis	13
3.2.1	Data Quality Check	13
3.2.2	Exploratory Data Analysis	14
3.2.3	Stationarity Testing	15
3.2.3.1	Augmented Dickey-Fuller Test	15
3.2.3.2	Kwiatkowski-Phillips-Schmidt-Shin Test	16
3.3	Data Pre-processing	16
3.3.1	Data Splitting using Sliding Windows	16
3.3.2	Data Scaling and Normalization	17
3.3.2.1	Z-Scale Normalization:	17
3.3.2.2	Min-Max Scaling:	18
3.3.3	Detrending and Addressing Non-Stationarity	18
3.3.4	Sequence Transformation	18
3.4	Model Building	19
3.4.1	LSTM Classifier	19
3.4.1.1	Structure	19
3.4.1.2	Training	19
3.4.1.2.1	Optimizer	19
3.4.1.2.1.1	Adam	19
3.4.1.2.1.2	Nadam	19
3.4.1.2.2	Loss Function	19
3.4.1.2.2.1	Binary Cross-entropy	19

3.4.1.2.2.2	Binary Focal Cross-entropy	19
3.4.1.3	Threshold Optimization	20
3.4.1.3.1	Fixed Thresholds	20
3.4.1.3.2	PR Curve	20
3.4.2	LSTM Regressor	20
3.4.2.1	Structure	20
3.4.2.2	Training	20
3.4.2.2.1	Optimizer	20
3.4.2.2.1.1	Adagrad	20
3.4.2.2.1.2	RMSprop	20
3.4.2.3	Loss Function	20
3.4.2.3.0.1	Mean Squared Error (mse)	20
3.4.3	Hybrid Models	20
3.4.3.1	GRU-LSTM	20
3.4.3.1.1	Structure	20
3.4.3.2	TCN-LSTM	20
3.4.3.2.1	Structure	20
3.5	Multivariate Models	21
3.5.1	Feature Selection	21
3.5.1.1	Spearman Correlation	21
3.5.1.2	Granger Causality Test	21
3.6	Hyperparameter Optimization	21
3.6.1	Grid Search	21
3.6.2	Random Search	21
3.6.3	Bayesian Optimization	21
3.6.4	Genetic Algorithm	21
3.6.5	Hyperband	21
3.6.6	Simulated Annealing	21
3.6.7	Reinforcement Learning	21

4	Evaluation Methods	22
4.1	Classification Metrics	22
4.2	Regression Metrics	22
4.3	Ablation Experiments	22
5	Development	23
6	Results and Discussion	24
7	Conclusion	25
7.1	Achievements	25
7.2	Limitations	25
7.3	Future Work	25

Abstract

插图

3.1	Line Chart of the USDEUR Exchange Rate (The Primary Dataset)	14
3.2	Rolling Mean and Standard Deviation of the USDEUR Exchange Rate (The Primary Dataset))	15
3.3	Sliding Window Data Splitting Strategy	17

表格

Chapter 1

Introduction

1.1 Context and Motivation

1.2 Aims and Objectives

1.3 Report Structure

Chapter 2

Background

2.1 The Foreign Exchange Market

2.2 Time Series Forecasting

2.2.1 Classification

2.2.2 Regression

2.3 Currency Exchange Rate Forecasting

2.3.1 Statistics Approach

2.3.2 Classical Machine Learning Approach

2.3.3 Deep Learning Approach

2.3.4 Hybrid Methods

2.4 Hyperparameter Optimization

2.5 Critical Analysis¹⁰

Chapter 3

Methodology

This chapter presents the methodology used to achieve the project’s objectives. It starts by describing the data collection process and analyzing the data collected, followed by an overview of the pre-processing techniques. The chapter then discusses the model building approach and the hyperparameter optimization strategy.

3.1 Data Collection

Data for this study were obtained from the Yahoo Finance platform using Python’s `yfinance` package. The data collection process was designed to support three main aspects of the research: (i) the primary dataset, (ii) supplementary features for multivariate modeling, and (iii) datasets for benchmarking against previous studies.

QUESTION: 我该不该引用yfinance的官方文档?

3.1.1 Primary Dataset

The primary dataset consists of the USDEUR exchange rate with a daily frequency, spanning from December 1, 2003 to January 31, 2025.

QUESTION: 我是否应该提及具体的数据收集的脚本

3.1.2 Supplementary Features for Multivariate Models

To build a robust multivariate model, additional financial indicators were collected. The supplementary data include:

- Crude Oil (WTI Futures)
- Gold Futures
- FTSE 100 Index
- US Dollar Index (DXY)

These datasets cover the period from January 1, 2000 to January 31, 2025 with a daily frequency. When used, they are aligned based on the corresponding currency pair's time base.

3.1.3 Benchmarking Datasets

For comparative analysis with prior research, additional datasets were collected to ensure that the time series forecasting results are directly comparable. Two sets of benchmarking data were collected:

1. A multi-currency dataset covering the period from December 18, 2017 to January 27, 2023. This dataset includes exchange rates for EUR/USD, GBP/USD, AUD/USD, and NZD/USD. For USD/JPY data, the script inverts the closing prices to derive the JPY/USD rate[1].

2. A focused subset for the EUR/USD pair spanning from January 1, 2013 to January 1, 2018[2].

QUESTION: 直接这样引用数据来源的论文，会不会有点突兀？ 写的很难受，该不该写Data Analysis? 总觉得不对劲

3.2 Data Analysis

Data analysis is a critical step in understanding the underlying quality, patterns, and characteristics of the dataset.

3.2.1 Data Quality Check

Ensuring high data quality is a critical first step before any analysis. In this stage, the following checks are performed:

- **Missing Values:** Confirm that there are no missing entries, or if there are, decide on an appropriate imputation method.
- **Outliers:** Identify any extreme values using statistical methods (this project used IQR), or if there are any, determine if they need to be removed or capped.
- **Duplicates:** Check for duplicate records to prevent bias in analysis.
- **Consistency and Integrity:** Ensure data types, ranges, and formats are consistent across the dataset.

Results: All datasets collected via yfinance contain no missing values or duplicates, and no outliers were detected using the Interquartile Range (IQR) method. Data types are consistent, with the ‘Date’ as an object and ‘Close’ as float64.

3.2.2 Exploratory Data Analysis

Exploratory Data Analysis (EDA) involves summarizing the main characteristics of the dataset using visual and quantitative methods. This project includes the following EDA techniques:

- **Summary Statistics:** Compute mean, median, variance, and other descriptive measures to understand data distribution.
- **Line Chart for Long-Term Trends:** Plot the USDEUR exchange rate over time to reveal long-term trends and identify potential anomalies. (See Figure 3.1.)
- **Rolling Statistics for Volatility Analysis:** Compute the rolling mean and standard deviation of the USDEUR exchange rate using a 30-day window to assess volatility changes over time. (See Figure 3.2.)




图 3.1: Line Chart of the USDEUR Exchange Rate (The Primary Dataset)




图 3.2: Rolling Mean and Standard Deviation of the USDEUR Exchange Rate (The Primary Dataset))

For the primary dataset, its summary statistics indicate that the mean exchange rate is approximately 0.823 with a standard deviation of 0.084, which reflects moderate variability. The line charts visually confirms this stability, which highlights long-term fluctuations without evident anomalies.

QUESTION: 加上非主要数据集的EDA会要写太多内容

3.2.3 Stationarity Testing

Stationarity testing is an important process in time series analysis. A stationary time series has statistical properties that do not change over time.

3.2.3.1 Augmented Dickey-Fuller Test

The Augmented Dickey-Fuller (ADF) test operates under the following principles:

- **Null Hypothesis (H_0):** The time series has a unit root (non-stationary).
- **Alternative Hypothesis (H_1):** The time series is stationary.

3.2.3.2 Kwiatkowski-Phillips-Schmidt-Shin Test

The Kwiatkowski-Phillips-Schmidt-Shin (KPSS) test operates under these principles:

- **Null Hypothesis (H_0):** The time series is stationary around a deterministic trend (level stationary).
- **Alternative Hypothesis (H_1):** The time series is non-stationary.

Results: The ADF test failed to reject the null hypothesis of a unit root at the 1% significance level, while the KPSS test rejected the null hypothesis of stationarity at the 1% level. Together, these results provide strong evidence that the dataset is non-stationary, which suggests that further differencing is required to achieve stationarity.

QUESTION: multivariate的feature selection部分不放在这里有影响吗?

3.3 Data Pre-processing

3.3.1 Data Splitting using Sliding Windows

For time series forecasting, model's performance can differ significantly over different time periods. To ensure the robustness and predictive reliability of my models across different time intervals, this project designs a sliding window approach for data splitting. Rather than segmenting the dataset directly, the sliding window method ensures effective use of overlapping intervals, which is particularly useful for short time series data.

The dataset is divided into training, validation, and test sets within each sliding window. The training set fits predictive models, the validation set is for hyperparameter tuning, and the test set provides an unbiased evaluation of the model's performance.

Figure 3.3 illustrates the sliding window approach, which shows how overlapping data intervals effectively split the dataset and maximize data usage.




图 3.3: Sliding Window Data Splitting Strategy

3.3.2 Data Scaling and Normalization

Data scaling and normalization are important preprocessing techniques aimed at improving model performance, convergence speed, and stability.

3.3.2.1 Z-Scale Normalization:

Z-scale normalization is defined as:

$$z = \frac{x - \mu}{\sigma} \quad (3.1)$$

where x is a data point, μ is the mean, and σ is the standard deviation of the dataset.

3.3.2.2 Min-Max Scaling:

Min-max scaling rescales data to the range $[0, 1]$:

$$x' = \frac{x - \min(x)}{\max(x) - \min(x)} \quad (3.2)$$

Min-max scaling is particularly useful when data distribution is not normal.

Note that it is vital to invert the scaling transformation so that the regression metrics are calculated on the original scale.

3.3.3 Detrending and Addressing Non-Stationarity

Time series data often exhibits trends that can negatively affect predictive modeling. Differencing is an effective technique to remove trends and improve the stationarity of time series data. Differencing involves computing the difference between consecutive observations:

$$x'_t = x_t - x_{t-1} \quad (3.3)$$

where x'_t is the differenced data at time t .

However, I failed to invert the differencing transformation for regression metrics calculation, thus the differencing was not applied to regression models. **QUESTION: 提及个人的错误是否合适?**

3.3.4 Sequence Transformation

Long Short-Term Memory (LSTM) models require data structured into sequences. This involves reshaping data into a three-dimensional array:

$$(\text{samples}, \text{timesteps}, \text{features}) \quad (3.4)$$

where samples represent the number of independent sequences in the dataset, timesteps represent the number of time intervals within each sample, and features represent the number of variables observed at each time step.

3.4 Model Building

This section introduces various deep learning approaches employed for building forecasting models on currency exchange rate. It discusses LSTM-based classifier and regressor, highlights their general structures and performance optimization techniques. Additionally, hybrid modeling methods are presented. Finally, multivariate modeling are briefly discussed.

3.4.1 LSTM Classifier

3.4.1.1 Structure

3.4.1.2 Training

3.4.1.2.1 Optimizer

3.4.1.2.1.1 Adam

3.4.1.2.1.2 Nadam

3.4.1.2.2 Loss Function

3.4.1.2.2.1 Binary Cross-entropy

3.4.1.2.2.2 Binary Focal Cross-entropy

3.4.1.3 Threshold Optimization

3.4.1.3.1 Fixed Thresholds

3.4.1.3.2 PR Curve

3.4.2 LSTM Regressor

3.4.2.1 Structure

3.4.2.2 Training

3.4.2.2.1 Optimizer

3.4.2.2.1.1 Adagrad

3.4.2.2.1.2 RMSprop

3.4.2.3 Loss Function

3.4.2.3.0.1 Mean Squared Error (mse)

3.4.3 Hybrid Models

3.4.3.1 GRU-LSTM

3.4.3.1.1 Structure

3.4.3.2 TCN-LSTM

3.4.3.2.1 Structure

3.5 Multivariate Models

3.5.1 Feature Selection

3.5.1.1 Spearman Correlation

3.5.1.2 Granger Causality Test

3.6 Hyperparameter Optimization

3.6.1 Grid Search

3.6.2 Random Search

3.6.3 Bayesian Optimization

3.6.4 Genetic Algorithm

3.6.5 Hyperband

3.6.6 Simulated Annealing

3.6.7 Reinforcement Learning

Chapter 4

Evaluation Methods

4.1 Classification Metrics

4.2 Regression Metrics

4.3 Ablation Experiments

Chapter 5

Development

Chapter 6

Results and Discussion

Chapter 7

Conclusion

7.1 Achievements

7.2 Limitations

7.3 Future Work

参考文献

- [1] Fernando García, Francisco Guijarro, Javier Oliver, and Rima Tamošiūnienė. Foreign exchange forecasting models: Arima and lstm comparison. *Engineering Proceedings*, 39(1):81, 2023.
- [2] D.C. Yıldırım, I.H. Toroslu, and U. Fiore. Forecasting directional movement of forex data using lstm with technical and macroeconomic indicators. *Financial Innovation*, 7(1), 2021.